

BDSE 15 第四組

Hadoop 大數據分析平台架設

聶廷宇

翁孟愉

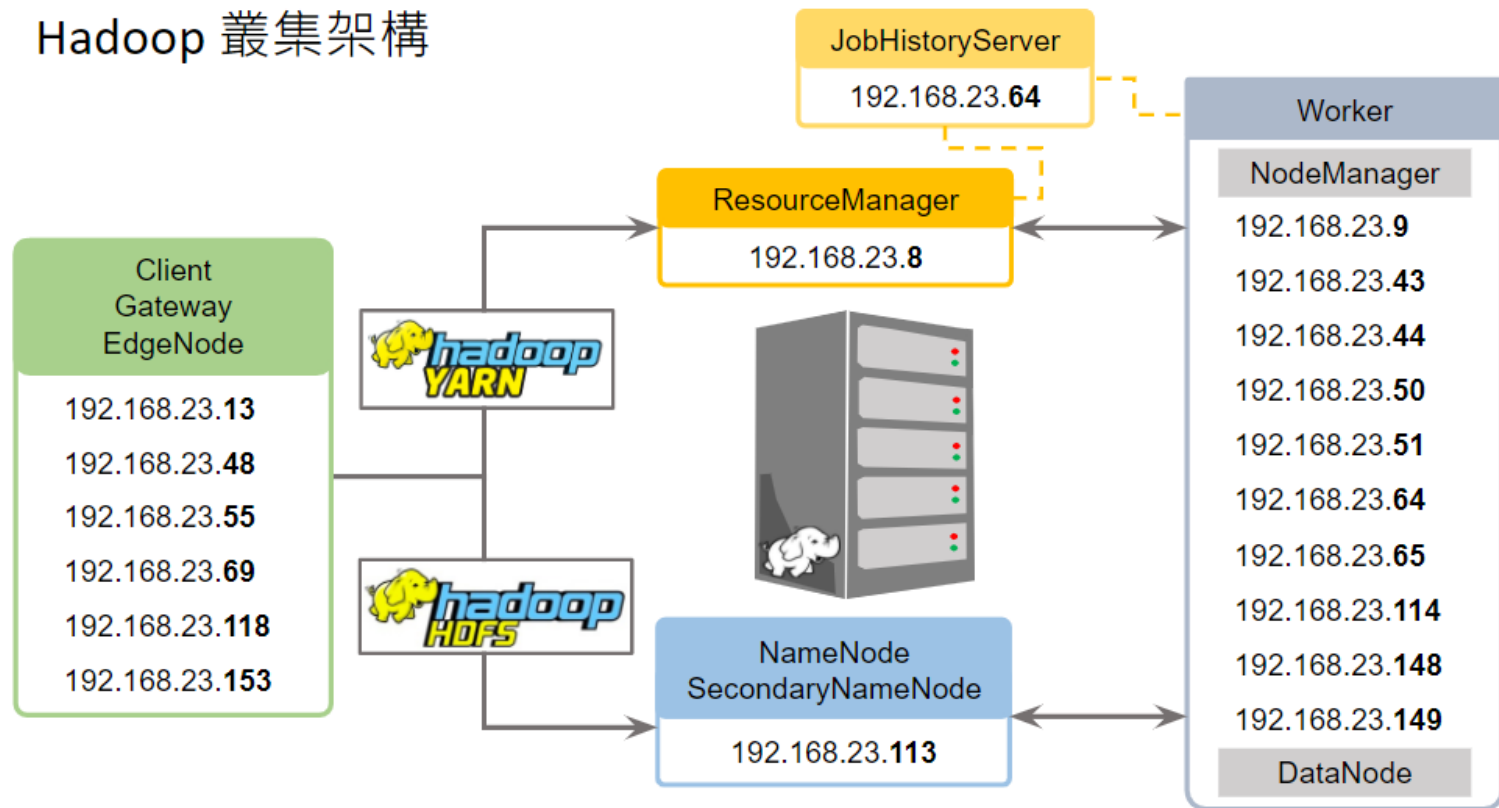
曾于芮

謝謙謙

林容聖

彭于靜

Hadoop 叢集架構



Hadoop叢集架構簡介

		節點名稱	節點功能
Hadoop Cluster	檔案系統： HDFS	NameNode (*1)	<ul style="list-style-type: none">負責管理和維護HDFS的名稱空間控制檔案的任何讀寫動作管理metadata
		Secondary NameNode (*1)	<ul style="list-style-type: none">定期在HDFS執行namespaces的checkpoint合併NameNode的edit logs到fsimage檔案，並寫入硬碟或檔案系統
		DataNode (*10)	<ul style="list-style-type: none">儲存真正的data
	運算框架： YARN	Resource Manager (*1)	<ul style="list-style-type: none">負責整個集群的資源分配
		Node Manager (*10)	<ul style="list-style-type: none">集群中的每個具體node的管理者負責執行與管理node上的containers
	Log record	jobHistory (*1)	<ul style="list-style-type: none">可通過web UI展示歷史作業日誌 (啟動時間、結束時間、各個任務的執行時間，各種Counter資料等)
Edge Nodes		Client (*6)	<ul style="list-style-type: none">Entry points to Hadoop clusterSparks jobs 撒入點

啟動Hadoop

```
2020-12-15 11:32:25,027 INFO common.Storage: Storage directory /home/hadoop/data/dfs/name has been successfully formatted.
```

NameNode電腦以hdfs namenode -format格式化的成功訊息

```
hadoop@bdsell13:~/data/dfs/name/current$ cat VERSION
#Tue Dec 15 11:32:24 CST 2020
namespaceID=1768217276
blockpoolID=BP-1241192631-192.168.23.113-1608003144930
storageType=NAME_NODE
cTime=1608003144930
clusterID=CID-38713cfe-5729-4e63-bd42-04b17062a2ec
layoutVersion=-65
```

NameNode電腦於格式化HDFS後可在
~/data/dfs/name/current查看VERSION以查詢clusterID

```
hadoop@bdsell13:~/data/dfs/name/current$ jps
6665 SecondaryNameNode
6957 Jps
6382 NameNode
```

HDFS開啟，NameNode電腦以jps指令確認NameNode已開啟

```
hadoop@bdsell14:/usr/local/hadoop/etc/hadoop$ jps
6117 Jps
5884 DataNode
```

HDFS開啟，Worker電腦以jps指令確認DataNode已開啟

```
hadoop@bdse8:/usr/local/hadoop/logs$ jps
1937 ResourceManager
8173 Jps
```

YARN開啟，ResourceManager電腦以jps指令確認ResourceManager已開啟

```
hadoop@bdsell14:/usr/local/hadoop/etc/hadoop$ jps
9488 NodeManager
9640 Jps
5884 DataNode
```

YARN開啟，Worker電腦以jps指令確認NodeManager已開啟

程式Demo

```
In [2]: # yarn mode
spark = SparkSession \
    .builder \
    .master("yarn") \
    .config('spark.executor.instances','9999') \
    .config('spark.executor.memory','2G') \
    .config('spark.executor.cores','1') \
    .appName("bdse118_Timmy") \
    .getOrCreate()
```

executor記憶體
仍有調整空間

```
In [7]: import pyspark
df = spark.read.csv("/user/hadoop/OpenSky/"+csv2019[0], header=True, inferSchema=True)
for i in csv2019[1:]:
    dfTemp = spark.read.csv("/user/hadoop/OpenSky/"+i, header=True, inferSchema=True)
    df = df.union(dfTemp)
```

```
In [8]: df.count(), len(df.columns)
```

```
Out[8]: (30989481, 16)
```

```
In [10]: %%time
df.groupby("origin").count().show()
```

```
In [11]: %%time
df.groupby("origin").count().show()
```

origin	count
LYBE	18160
YTQY	733
3AK4	44
KPKK	22660
KLCK	2614
RPUF	86
SDYJ	588
KAVP	1586
LQWK	1023
WN61	994
K08A	462
VA16	1083
4FL9	1177
FD49	166
EKRA	39
CL35	375
41NC	150
LOGP	102
LE1B	2287
1GE5	5

only showing top 20 rows

```
CPU times: user 0 ns, sys: 4.04 ms, total: 4.04 ms
Wall time: 10.1 s
```

Spark on YARN

	origin	count
進行比較	YTQY	733
	VA16	1083
	4FL9J	1177
	SDYJ	588
	FD49	166
	CL35	375
	KAVP	1586
	KPKD	22660
	KFSI	12
	5M03	73
	KAXS	64
	WN61	994
	8LL0	47
	SJPH	18
	ESSX	514
	KE41	75
K6G1	94	
PAWD	49	
K08A	462	
NZWK	8	

only showing top 20 rows

Spark Local
執行24秒

進行比較

FinishTime	State	FinalStatus	Running Containers	Allocated CPU VCores	Allocated Memory MB	Reserved CPU VCores	Reserved Memory MB	% of Queue	% of Cluster
N/A	RUNNING	UNDEFINED	61	61	185344	6	18432	90.5	90.5

Cluster

- Home
- Nodes
- Jobs
- Submits
- Applications
- Events
- Configuration
- Log Explorer
- Alerts
- Help

Tools

- Dashboard
- Jobs
- Nodes
- Submits
- Applications
- Events
- Configuration
- Log Explorer
- Alerts
- Help

Application Attempt appattemp_attempt_1608081189355_0039_000001

Application Attempt Status: RUNNING
 Time: Wed Jul 12 22:42 +0800 2009
Framework: ChRM, YSAC
Executor: container, NodeContainer955, 9639, 01, 000001
AM Container: Node: N/A
 Tracking URL: nodeattempattempt
Diagnosis Info: Nodes blacklisted by the application: -
 Nodes blacklisted by the system: -

[Application Attempt Metrics](#)

Application Attempt Headers: <memory 1024, <cores 4>

Each table cell represents the number of Nodes, or RackLocal or OffSwitch containers by NodeId, or RackLocal or OffSwitch resource requests

	Node Local Request	Rack Local Request	Off Switch Request
Num Node Local Containers (satisfied by)	0		
Num Rack Local Containers (satisfied by)	0		
Num Off Switch Containers (satisfied by)	0	0	61

Total Outstanding Resource Requests: <memory 36532408, <cores 939>

View 20 [Resources](#)

Priority	Application/RequestId	Resource/Name	Capability	Num-Containers	Relative-Localty	Node-Label/Expression	Execution-Time	Allocation-Status	Placement-Constraint
1	*	memory 3672, <cores 1>	939B	0	true		GUARANTEED	N/A	N/A

Showing 1 of 1 entries

[Expansions in cache](#) [For more details view Job Applications](#) [Schedule Activities](#) [Refresh](#)

Logged in as [jdoe](#)

Container ID	Node	Container Exit Status	Logs
container_1608081189355_0039_01_000001	http://bde1416.example.org:3042		Logs
container_1608081189355_0039_01_000000	http://bde1416.example.org:3042		Logs
container_1608081189355_0039_01_000009	http://bde1416.example.org:3042		Logs
container_1608081189355_0039_01_000008	http://bde1416.example.org:3042		Logs
container_1608081189355_0039_01_000007	http://bde1416.example.org:3042		Logs
container_1608081189355_0039_01_000006	http://bde1416.example.org:3042		Logs
container_1608081189355_0039_01_000005	http://bde1416.example.org:3042		Logs
container_1608081189355_0039_01_000004	http://bde1416.example.org:3042		Logs
container_1608081189355_0039_01_000003	http://bde1416.example.org:3042		Logs
container_1608081189355_0039_01_000002	http://bde1416.example.org:3042		Logs

執行程式時，同步檢視
container數量、記憶體
用量以利後續調整

故障排除

Troubleshooting



執行情境：
於 Spark YARN 執行程式



遭遇問題：
無法識別bdse8.example.org

```
hadoop@bdse55: /usr/local/spark
resa@bdse55:~$ sudo -i
[sudo] password for resa:
root@bdse55:~# cp ~hadoop/hosts /etc/hosts -l /etc/hosts
cp: target '/etc/hosts' is not a directory
root@bdse55:~# exit
logout
resa@bdse55:~$ su - hadoop
Password:
hadoop@bdse55:~$ cd $SPARK_HOME
hadoop@bdse55:~/usr/local/spark$ bin/spark-submit --class org.apache.spark.examples.SparkPi \
> --master yarn
> --deploy-mode
> --driver-memory
> --executor-memory
> --executor-cpu
> --num-executor
> examples/jars/
> 10
WARNING: An illegal reflective operation was detected
WARNING: Illegal reflective operation by java.lang.ProcessImpl
r/local/spark/jars/s
teBuffer(long,int)
WARNING: Please consider reporting this to the vendor of the
nsafe.Platform
WARNING: Use --illegal-access=warn to enable warnings of this type.
ive access operation
WARNING: All illegal access operations will be denied in a future release
2020-12-15 17:00:25,
library for your pl
2020-12-15 17:00:26,
se8.example.org:8032
2020-12-15 17:00:26,
hadoop@bdse55:~/usr/local/spark$
hadoop@bdse55:~/usr/local/spark$
invoking ApplicationClientProtocolPBClientImpl.getClusterMetrics over null after 10 failover at
tempts. Trying to failover after sleeping for 18808ms.
2020-12-15 17:05:58,744 INFO retry.RetryInvocationHandler: java.net.UnknownHostException: Inval
id host name: local host is: (unknown); destination host is: "bdse8.example.org":8032; java.net
.UnknownHostException: For more details see: http://wiki.apache.org/hadoop/UnknownHost, while
invoking ApplicationClientProtocolPBClientImpl.getClusterMetrics over null after 11 failover at
tempts. Trying to failover after sleeping for 31651ms.
2020-12-15 17:06:30,396 INFO retry.RetryInvocationHandler: java.net.UnknownHostException: Inval
id host name: local host is: (unknown); destination host is: "bdse8.example.org":8032; java.net
.UnknownHostException: For more details see: http://wiki.apache.org/hadoop/UnknownHost, while
invoking ApplicationClientProtocolPBClientImpl.getClusterMetrics over null after 12 failover at
tempts. Trying to failover after sleeping for 22577ms.
2020-12-15 17:06:52,974 INFO retry.RetryInvocationHandler: java.net.UnknownHostException: Inval
id host name: local host is: (unknown); destination host is: "bdse8.example.org":8032; java.net
.UnknownHostException: For more details see: http://wiki.apache.org/hadoop/UnknownHost, while
invoking ApplicationClientProtocolPBClientImpl.getClusterMetrics over null after 13 failover at
tempts. Trying to failover after sleeping for 15033ms.
2020-12-15 17:07:08,008 INFO retry.RetryInvocationHandler: java.net.UnknownHostException: Inval
id host name: local host is: (unknown); destination host is: "bdse8.example.org":8032; java.net
.UnknownHostException: For more details see: http://wiki.apache.org/hadoop/UnknownHost, while
invoking ApplicationClientProtocolPBClientImpl.getClusterMetrics over null after 14 failover at
tempts. Trying to failover after sleeping for 26596ms.
2020-12-15 17:07:34,619 INFO retry.RetryInvocationHandler: java.net.UnknownHostException: Inval
id host name: local host is: (unknown); destination host is: "bdse8.example.org":8032; java.net
.UnknownHostException: For more details see: http://wiki.apache.org/hadoop/UnknownHost, while
invoking ApplicationClientProtocolPBClientImpl.getClusterMetrics over null after 15 failover at
tempts. Trying to failover after sleeping for 24024ms.
2020-12-15 17:07:58,645 INFO retry.RetryInvocationHandler: java.net.UnknownHostException: Inval
id host name: local host is: (unknown); destination host is: "bdse8.example.org":8032; java.net.Unkno
wnHostException: For more details see: http://wiki.apache.org/hadoop/UnknownHost, while invok
ing ApplicationClientProtocolPBClientImpl.getClusterMetrics over null after 16 failover attempts. Tryi
ng to failover after sleeping for 27794ms.
```

故障排除

Troubleshooting



原因偵查：

1. 檢查 yarn-site.xml 檔案內容
-> 設定正確
2. 改檢查 /etc/hosts
-> 發現沒有把 bdse8 加入檔案內



解決方法：

重新編輯 /etc/hosts 檔案內容
加入正確 cluster FQDN



修正後執行分析程式成功！

yarn-site.xml

```
hadoop@bdse55:/usr/local/hadoop/etc/hadoop$ tail -5 yarn
tail: cannot open 'yarn' for reading: No such file or directory
hadoop@bdse55:/usr/local/hadoop/etc/hadoop$ tail -5 yarn-site.xml
<property>
  <name>yarn.resourcemanager.hostname</name>
  <value>bdse8.example.org</value>
</property>
</configuration>
```

檔案內設定無誤

/etc/hosts

```
192.168.23.13 bdse13.example.org bdse13
192.168.23.48 bdse48.example.org bdse48
192.168.23.55 bdse55.example.org bdse55
192.168.23.69 bdse69.example.org bdse69
192.168.23.118 bdse118.example.org bdse118
192.168.23.153 bdse153.example.org bdse153
```

無bdse8

/etc/hosts

```
GNU nano 4.8 /etc/hosts
1 27.0.0.1 localhost
# Momo
192.168.23.8 bdse8.example.org bdse8
192.168.23.9 bdse9.example.org bdse9
192.168.23.13 bdse13.example.org bdse13
# Ula
192.168.23.43 bdse43.example.org bdse43
192.168.23.44 bdse44.example.org bdse44
192.168.23.48 bdse48.example.org bdse48
# Yujui
192.168.23.50 bdse50.example.org bdse50
192.168.23.51 bdse51.example.org bdse51
192.168.23.55 bdse55.example.org bdse55
```

補上bdse8

故障排除

Troubleshooting



問題原因：

為提升叢集效能，將Spark程式庫統一上傳到HDFS。
但在執行分析程式(Spark on YARN mode)時，出現Exception。



解決方法：

1. 從Exception發現，Gateway接到自己的電腦，但正確應該是要接到NameNode那台電腦。
2. 在設定檔spark-defaults.conf中，找到有關叢集效能提升的程式碼，並做修正。

Exception

```
/usr/local/spark
...
2020-12-16 10:23:36,537 INFO yarn.Client: Will allocate AM container, with 1408
MB memory including 384 MB overhead
2020-12-16 10:23:36,537 INFO yarn.Client: Setting up container launch context f
r our AM
2020-12-16 10:23:36,538 INFO yarn.Client: Setting up the launch environment for
our AM container
2020-12-16 10:23:36,547 INFO yarn.Client: Preparing resources for our AM conta
iner
2020-12-16 10:23:36,616 INFO yarn.Client: Deleted staging directory hdfs://bdse
13.example.org/user/hadoop/_sparkStaging/application_1608081189355_0020
Exception in thread "main" java.net.ConnectException: Call From bdse69/192.168.
3.69 to bdse69.example.org:8020 failed on connection exception: java.net.Connec
Exception: Connection refused; For more details see: http://wiki.apache.org/ha
adoop/ConnectionRefused
...
at java.base
nce0 (Native Method)
at java.base
nce (NativeConstruct
instance (DelegatingCo
490)
at java.base
at org.apach
```

發現問題並解決

```
spark-defaults.conf Modified
# Example:
# spark.master spark://master:7077
# spark.eventLog.enabled true
# spark.eventLog.dir hdfs://namenode:8021/directory
# spark.serializer org.apache.spark.serializer.KryoSerializer
# spark.driver.memory 5g
# spark.executor.extraJavaOptions -XX:+PrintGCDetails -Dkey=value -Dnumbers="1 2 3"
spark.driver.extraJavaOptions="-Dio.netty.tryReflectionSetAccessible=true"
spark.executor.extraJavaOptions="-Dio.netty.tryReflectionSetAccessible=true"
spark.yarn.jars hdfs://bdse69.example.org/user/spark/share/jars/*
```

改成bdse113(本組NameNode)

Thanks!